



limhpone / computervision-final-prep



Code

Issues

Pull requests

Actions

Projects

Wiki

Security

In

computervision-final-prep / lab / Lab 08 (YOLO)-20251128 / utils.py



limhpone YOLO

fd90b91 · 2 hours ago



349 lines (278 loc) · 11.6 KB

Code

Blame



Raw



```
1     import torch
2     import numpy as np
3     import matplotlib.pyplot as plt
4     import matplotlib.patches as patches
5     from collections import Counter
6
7     def intersection_over_union(boxes_preds, boxes_labels, box_format="midpoint"):
8         """
9             Calculates intersection over union
10
11            Parameters:
12                boxes_preds (tensor): Predictions of Bounding Boxes (BATCH_SIZE, 4)
13                boxes_labels (tensor): Correct labels of Bounding Boxes (BATCH_SIZE, 4)
14                box_format (str): midpoint/corners, if boxes (x,y,w,h) or (x1,y1,x2,y2)
15
16            Returns:
17                tensor: Intersection over union for all examples
18        """
19
20        if box_format == "midpoint":
21            box1_x1 = boxes_preds[:, 0:1] - boxes_preds[:, 2:3] / 2
22            box1_y1 = boxes_preds[:, 1:2] - boxes_preds[:, 3:4] / 2
23            box1_x2 = boxes_preds[:, 0:1] + boxes_preds[:, 2:3] / 2
24            box1_y2 = boxes_preds[:, 1:2] + boxes_preds[:, 3:4] / 2
25            box2_x1 = boxes_labels[:, 0:1] - boxes_labels[:, 2:3] / 2
26            box2_y1 = boxes_labels[:, 1:2] - boxes_labels[:, 3:4] / 2
27            box2_x2 = boxes_labels[:, 0:1] + boxes_labels[:, 2:3] / 2
28            box2_y2 = boxes_labels[:, 1:2] + boxes_labels[:, 3:4] / 2
29
30        if box_format == "corners":
31            box1_x1 = boxes_preds[:, 0:1]
32            box1_y1 = boxes_preds[:, 1:2]
33            box1_x2 = boxes_preds[:, 2:3]
```

```
34     box1_y2 = boxes_preds[..., 3:4] # (N, 1)
35     box2_x1 = boxes_labels[..., 0:1]
36     box2_y1 = boxes_labels[..., 1:2]
37     box2_x2 = boxes_labels[..., 2:3]
38     box2_y2 = boxes_labels[..., 3:4]
39
40     x1 = torch.max(box1_x1, box2_x1)
41     y1 = torch.max(box1_y1, box2_y1)
42     x2 = torch.min(box1_x2, box2_x2)
43     y2 = torch.min(box1_y2, box2_y2)
44
45     # .clamp(0) is for the case when they do not intersect
46     intersection = (x2 - x1).clamp(0) * (y2 - y1).clamp(0)
47
48     box1_area = abs((box1_x2 - box1_x1) * (box1_y2 - box1_y1))
49     box2_area = abs((box2_x2 - box2_x1) * (box2_y2 - box2_y1))
50
51     return intersection / (box1_area + box2_area - intersection + 1e-6)
52
53
54     def non_max_suppression(bboxes, iou_threshold, threshold, box_format="corners"):
55         """
56             Does Non Max Suppression given bboxes
57
58             Parameters:
59                 bboxes (list): list of lists containing all bboxes with each bboxes
59                 specified as [class_pred, prob_score, x1, y1, x2, y2]
60                 iou_threshold (float): threshold where predicted bboxes is correct
61                 threshold (float): threshold to remove predicted bboxes (independent of IoU)
62                 box_format (str): "midpoint" or "corners" used to specify bboxes
63
64             Returns:
65                 list: bboxes after performing NMS given a specific IoU threshold
66
66             """
67
68
69     assert type(bboxes) == list
70
71     bboxes = [box for box in bboxes if box[1] > threshold]
72     bboxes = sorted(bboxes, key=lambda x: x[1], reverse=True)
73     bboxes_after_nms = []
74
75     while bboxes:
76         chosen_box = bboxes.pop(0)
77
78         bboxes = [
79             box
80             for box in bboxes
81             if box[0] != chosen_box[0]
82             or intersection_over_union(

```

```
83         torch.tensor(chosen_box[2:]),
84         torch.tensor(box[2:]),
85         box_format=box_format,
86     )
87     < iou_threshold
88 ]
89
90     bboxes_after_nms.append(chosen_box)
91
92     return bboxes_after_nms
93
94
95     def mean_average_precision(
96         pred_boxes, true_boxes, iou_threshold=0.5, box_format="midpoint", num_classes=20
97     ):
98         """
99             Calculates mean average precision
100
101         Parameters:
102             pred_boxes (list): list of lists containing all bboxes with each bboxes
103             specified as [train_idx, class_prediction, prob_score, x1, y1, x2, y2]
104             true_boxes (list): Similar as pred_boxes except all the correct ones
105             iou_threshold (float): threshold where predicted bboxes is correct
106             box_format (str): "midpoint" or "corners" used to specify bboxes
107             num_classes (int): number of classes
108
109         Returns:
110             float: mAP value across all classes given a specific IoU threshold
111         """
112
113         # list storing all AP for respective classes
114         average_precisions = []
115
116         # used for numerical stability later on
117         epsilon = 1e-6
118
119         for c in range(num_classes):
120             detections = []
121             ground_truths = []
122
123             # Go through all predictions and targets,
124             # and only add the ones that belong to the
125             # current class c
126             for detection in pred_boxes:
127                 if detection[1] == c:
128                     detections.append(detection)
129
130                 for true_box in true_boxes:
131                     if true_box[1] == c:
```

```
132         ground_truths.append(true_box)
133
134     # find the amount of bboxes for each training example
135     # Counter here finds how many ground truth bboxes we get
136     # for each training example, so let's say img 0 has 3,
137     # img 1 has 5 then we will obtain a dictionary with:
138     # amount_bboxes = {0:3, 1:5}
139     amount_bboxes = Counter([gt[0] for gt in ground_truths])
140
141     # We then go through each key, val in this dictionary
142     # and convert to the following (w.r.t same example):
143     # amount_bboxes = {0:torch.tensor[0,0,0], 1:torch.tensor[0,0,0,0,0]}
144     for key, val in amount_bboxes.items():
145         amount_bboxes[key] = torch.zeros(val)
146
147     # sort by box probabilities which is index 2
148     detections.sort(key=lambda x: x[2], reverse=True)
149     TP = torch.zeros((len(detections)))
150     FP = torch.zeros((len(detections)))
151     total_true_bboxes = len(ground_truths)
152
153     # If none exists for this class then we can safely skip
154     if total_true_bboxes == 0:
155         continue
156
157     for detection_idx, detection in enumerate(detections):
158         # Only take out the ground_truths that have the same
159         # training idx as detection
160         ground_truth_img = [
161             bbox for bbox in ground_truths if bbox[0] == detection[0]
162         ]
163
164         num_gts = len(ground_truth_img)
165         best_iou = 0
166
167         for idx, gt in enumerate(ground_truth_img):
168             iou = intersection_over_union(
169                 torch.tensor(detection[3:]),
170                 torch.tensor(gt[3:]),
171                 box_format=box_format,
172             )
173
174             if iou > best_iou:
175                 best_iou = iou
176                 best_gt_idx = idx
177
178             if best_iou > iou_threshold:
179                 # only detect ground truth detection once
180                 if amount_bboxes[detection[0]][best_gt_idx] == 0:
```

```
181             # true positive and add this bounding box to seen
182             TP[detection_idx] = 1
183             amount_bboxes[detection[0]][best_gt_idx] = 1
184         else:
185             FP[detection_idx] = 1
186
187         # if IOU is lower then the detection is a false positive
188     else:
189         FP[detection_idx] = 1
190
191     TP_cumsum = torch.cumsum(TP, dim=0)
192     FP_cumsum = torch.cumsum(FP, dim=0)
193     recalls = TP_cumsum / (total_true_bboxes + epsilon)
194     precisions = torch.divide(TP_cumsum, (TP_cumsum + FP_cumsum + epsilon))
195     precisions = torch.cat((torch.tensor([1]), precisions))
196     recalls = torch.cat((torch.tensor([0]), recalls))
197     # torch.trapz for numerical integration
198     average_precisions.append(torch.trapz(precisions, recalls))
199
200     return sum(average_precisions) / len(average_precisions)
201
202
203     def plot_image(image, boxes):
204         """Plots predicted bounding boxes on the image"""
205         im = np.array(image)
206         height, width, _ = im.shape
207
208         # Create figure and axes
209         fig, ax = plt.subplots(1)
210         # Display the image
211         ax.imshow(im)
212
213         # box[0] is x midpoint, box[2] is width
214         # box[1] is y midpoint, box[3] is height
215
216         # Create a Rectangle patch
217         for box in boxes:
218             box = box[2:]
219             assert len(box) == 4, "Got more values than in x, y, w, h, in a box!"
220             upper_left_x = box[0] - box[2] / 2
221             upper_left_y = box[1] - box[3] / 2
222             rect = patches.Rectangle(
223                 (upper_left_x * width, upper_left_y * height),
224                 box[2] * width,
225                 box[3] * height,
226                 linewidth=1,
227                 edgecolor="r",
228                 facecolor="none",
229             )
```

```
230         # Add the patch to the Axes
231         ax.add_patch(rect)
232
233     plt.show()
234
235     def get_bboxes(
236         loader,
237         model,
238         iou_threshold,
239         threshold,
240         pred_format="cells",
241         box_format="midpoint",
242         device="cuda",
243     ):
244         all_pred_boxes = []
245         all_true_boxes = []
246
247         # make sure model is in eval before get bboxes
248         model.eval()
249         train_idx = 0
250
251         for batch_idx, (x, labels) in enumerate(loader):
252             x = x.to(device)
253             labels = labels.to(device)
254
255             with torch.no_grad():
256                 predictions = model(x)
257
258             batch_size = x.shape[0]
259             true_bboxes = cellboxes_to_boxes(labels)
260             bboxes = cellboxes_to_boxes(predictions)
261
262             for idx in range(batch_size):
263                 nms_boxes = non_max_suppression(
264                     bboxes[idx],
265                     iou_threshold=iou_threshold,
266                     threshold=threshold,
267                     box_format=box_format,
268                 )
269
270
271             #if batch_idx == 0 and idx == 0:
272             #    plot_image(x[idx].permute(1,2,0).to("cpu"), nms_boxes)
273             #    print(nms_boxes)
274
275             for nms_box in nms_boxes:
276                 all_pred_boxes.append([train_idx] + nms_box)
277
278             for box in true_bboxes[idx]:
```

```
279         # many will get converted to 0 pred
280         if box[1] > threshold:
281             all_true_boxes.append([train_idx] + box)
282
283             train_idx += 1
284
285         model.train()
286
287     return all_pred_boxes, all_true_boxes
288
289
290     def convert_cellboxes(predictions, S=7):
291         """
292             Converts bounding boxes output from Yolo with
293             an image split size of S into entire image ratios
294             rather than relative to cell ratios. Tried to do this
295             vectorized, but this resulted in quite difficult to read
296             code... Use as a black box? Or implement a more intuitive,
297             using 2 for loops iterating range(S) and convert them one
298             by one, resulting in a slower but more readable implementation.
299         """
300
301     predictions = predictions.to("cpu")
302     batch_size = predictions.shape[0]
303     predictions = predictions.reshape(batch_size, 7, 7, 30)
304     bboxes1 = predictions[..., 21:25]
305     bboxes2 = predictions[..., 26:30]
306     scores = torch.cat(
307         (predictions[..., 20].unsqueeze(0), predictions[..., 25].unsqueeze(0)), dim=0
308     )
309     best_box = scores.argmax(0).unsqueeze(-1)
310     best_boxes = bboxes1 * (1 - best_box) + best_box * bboxes2
311     cell_indices = torch.arange(7).repeat(batch_size, 7, 1).unsqueeze(-1)
312     x = 1 / S * (best_boxes[..., :1] + cell_indices)
313     y = 1 / S * (best_boxes[..., 1:2] + cell_indices.permute(0, 2, 1, 3))
314     w_y = 1 / S * best_boxes[..., 2:4]
315     converted_bboxes = torch.cat((x, y, w_y), dim=-1)
316     predicted_class = predictions[..., :20].argmax(-1).unsqueeze(-1)
317     best_confidence = torch.max(predictions[..., 20], predictions[..., 25]).unsqueeze(
318         -1
319     )
320     converted_preds = torch.cat(
321         (predicted_class, best_confidence, converted_bboxes), dim=-1
322     )
323
324     return converted_preds
325
326
327     def cellboxes_to_boxes(out, S=7):
```

```
328     converted_pred = convert_cellboxes(out).reshape(out.shape[0], S * S, -1)
329     converted_pred[..., 0] = converted_pred[..., 0].long()
330     all_bboxes = []
331
332     for ex_idx in range(out.shape[0]):
333         bboxes = []
334
335         for bbox_idx in range(S * S):
336             bboxes.append([x.item() for x in converted_pred[ex_idx, bbox_idx, :]])
337         all_bboxes.append(bboxes)
338
339     return all_bboxes
340
341 def save_checkpoint(state, filename="my_checkpoint.pth.tar"):
342     print("=> Saving checkpoint")
343     torch.save(state, filename)
344
345
346     def load_checkpoint(checkpoint, model, optimizer):
347         print("=> Loading checkpoint")
348         model.load_state_dict(checkpoint["state_dict"])
349         optimizer.load_state_dict(checkpoint["optimizer"])
```