

11/28/25, 7:57 AM

computervision-final-prep/lab/Lab 12 (3D Vision)-20251128/lab12.1_Pytorch3D_Rendering_Tutorial.ipynb at main · limhpone/com...



limhpone / computervision-final-prep



Code

Issues

Pull requests

Actions

Projects

Wiki

Security

In

computervision-final-prep / lab / Lab 12 (3D Vision)-20251128

/ lab12.1_Pytorch3D_Rendering_Tutorial.ipynb



limhpone 3D Vision- lab 12

daadd1a · 2 hours ago



2.62 MB



Pytorch3D Tutorial

Contents of tutorial:

- Installing Pytorch3D on Colab
- Mesh Rendering
- Point cloud Rendering

More tutorials available: <https://github.com/facebookresearch/pytorch3d>

Installing Pytorch3D

It will be much faster if you switch to GPU runtime (Runtime > Change runtime type > Hardware Accelerator > GPU).

In [83]:

```
import os
import sys
import torch

# Suppress scientific notation
torch.set_printoptions(sci_mode=False)

# Install PyTorch3D if necessary
need_pytorch3d=False

try:
    import pytorch3d
except ModuleNotFoundError:
    need_pytorch3d=True

if need_pytorch3d:
    !uv pip install --extra-index-url https://miopsota.github.io/torch_pac
```



In [84]:

```
!pwd
```

```
/content
```

In [85]:

```
import cv2
import matplotlib.pyplot as plt
import numpy as np
import pytorch3d
import pytorch3d.io
from pytorch3d.vis.plotly_vis import plot_scene
from tqdm.auto import tqdm
```

In [86]:

```
# This should print True if you are using your GPU
print("Using GPU:" torch.cuda.is_available())
```

```

print("Using GPU: ", torch.cuda.is_available())
if torch.cuda.is_available():
    device = torch.device("cuda")
else:
    device = torch.device("cpu")

```

Using GPU: True

Rendering a Mesh

To render a mesh, we need 3 components:

- A mesh to render
- A camera to render from
- A renderer (consisting of a rasterizer and shader)

Setting up a mesh

Take an example object data/cow.obj

In this file you will see 4 data types:

1. mtllib cow.mtl : This tells OBJ which MTL file to load.
2. v 0.10 0.20 0.30 : Lines starting with v hold vertex positions.
3. vt 0.50 0.75 : Lines starting with vt hold texture coordinates.
4. f 1/1 2/2 3/3 : Lines starting with f hold faces.

Each face entry gives vertex, texture index, and optional normal index.

```

faces.verts_idx → vertex indices
faces.textures_idx → uv indices

```

In [87]:

```

vertices, face_props, text_props = pytorch3d.io.load_obj("data/cow.obj")
faces = face_props.verts_idx

```

In [88]:

```

print("Vertices", vertices.shape)
print("Faces", faces.shape)

```

```

Vertices torch.Size([2930, 3])
Faces torch.Size([5856, 3])

```

In [89]:

```

print("vertices: ", vertices[0])
print("faces: ", faces[0])

faces.min(), faces.max()

```

```

vertices: tensor([ 0.3488, -0.3350, -0.0832])
faces: tensor([738, 734, 735])

```

Out[89]: (tensor(0), tensor(2929))

In [90]: # All Pytorch3D elements need to be hatched!

```
computervision-final-prep/lab/Lab 12 (3D Vision)-20251128/lab12.1_Pytorch3D_Rendering_Tutorial.ipynb at main · limhpon/com...
```

```
" All pytorch3d elements need to be batches."
vertices = vertices.unsqueeze(0) # 1 x num_vertices x 3
faces = faces.unsqueeze(0)       # 1 x num_faces x 3

print("Vertices", vertices.shape)
print("Faces", faces.shape)
```

```
Vertices torch.Size([1, 2930, 3])
Faces torch.Size([1, 5856, 3])
```

In [91]:

```
texture_rgb = torch.ones_like(vertices) # N X 3
texture_rgb = texture_rgb * torch.tensor([0.7, 0.7, 1])

# Wrap the colors as vertex textures
textures = pytorch3d.renderer.TexturesVertex(texture_rgb) # Manually create
print(texture_rgb)
print(texture_rgb.shape)

tensor([[0.7000, 0.7000, 1.0000],
        [0.7000, 0.7000, 1.0000],
        [0.7000, 0.7000, 1.0000],
        ...,
        [0.7000, 0.7000, 1.0000],
        [0.7000, 0.7000, 1.0000],
        [0.7000, 0.7000, 1.0000]]])
torch.Size([1, 2930, 3])
```

In [92]:

```
meshes = pytorch3d.structures.Meshes(
    verts=vertices, # batched tensor or a list of tensors
    faces=faces,
    textures=textures,
)
meshes = meshes.to(device) # Move mesh to GPU
```

In [93]:

```
print(meshes.verts_padded().shape)
print(meshes.faces_padded().shape)

torch.Size([1, 2930, 3])
torch.Size([1, 5856, 3])
```

Setting up a camera

In [94]:

```
R = torch.eye(3).unsqueeze(0) # identity rotation
T = torch.tensor([[0, 0, 3]]) # camera at z=3
```

In [95]:

```
# world -> view transforms
# view_T_world @ world = view

cameras = pytorch3d.renderer.FoVPerspectiveCameras(
    R=R,
    T=T,
    fov=60,
    device=device,
```

```
'  

# FoVPerspectiveCameras uses R and T: view_point = R * world_point + T
```

In [96]: cameras.get_camera_center() # returns the camera location in world space.
Since we have identity rotation and T moves the whole world by +3 in came

Out[96]: tensor([[0., 0., -3.]], device='cuda:0')

In [97]: transform = cameras.get_world_to_view_transform() # gives the final world
transform.get_matrix()

Out[97]: tensor([[[1., 0., 0., 0.],
[0., 1., 0., 0.],
[0., 0., 1., 0.],
[0., 0., 3., 1.]]], device='cuda:0')

This shifts the world by +3 along z in view space, which is the same as placing the camera at z = -3.

Setting up a Renderer

Rasterizer: Given a pixel, which triangles correspond to it?

Shader: Given triangle, texture, lighting, etc, how should the pixel be colored?

In [98]: image_size = 512
raster_settings = pytorch3d.renderer.RasterizationSettings(image_size=image_size)
rasterizer = pytorch3d.renderer.MeshRasterizer(
 raster_settings=raster_settings,
)
shader = pytorch3d.renderer.HardPhongShader(device=device) # applies light
renderer = pytorch3d.renderer.MeshRenderer(
 rasterizer=rasterizer,
 shader=shader,
) # a wrapper that first calls rasterizer(meshes, cameras) and shader (Lig


Render an image

In [99]: image = renderer(meshes, cameras=cameras)
print(image.shape)
torch.Size([1, 512, 512, 4])

In [100...]: image = image[0].cpu().numpy()

```
plt.imshow(image)
```

Out[100... <matplotlib.image.AxesImage at 0x7f0fd638f830>



In [101...


```
print(image.min(), image.max())
```


0.0 1.0

In [102... *# Lets add some lights to our image*

```
lights = pytorch3d.renderer.PointLights(location=[[0, 0, -3]], device=device)
image = renderer(meshes, cameras=cameras, lights=lights)
plt.imshow(image[0].cpu().numpy())
```

Out[102... <matplotlib.image.AxesImage at 0x7f0fd5c8f8f0>





In [103...]

```
plot_scene({
    "figure": {
        "Mesh": meshes,
        "Camera": cameras,
    }
})  

# plot_scene does not show the rendered image from the shader. It shows the
```



Common Rendering Problems

- Nothing is visible
 - Check Camera center
 - Check the Camera world_to_view transform. Try applying the transform to the origin or specific vertices of the mesh.
 - Use the Plotly visualization.
- Lights doesn't seem to do anything
 - Make sure there's no typo (must be `renderer(..., lights=lights)` and not `light=lights`)
 - Check the position of the light

Transformations

2 ways to apply geometric operations:

- Move the mesh
- Move the camera

In [104...]

```
relative_rotation = pytorch3d.transforms.euler_angles_to_matrix(
    torch.tensor([0, np.pi/2, 0]), "XYZ"
)
relative_rotation # rotation of 90 degrees around the Y axis
```

Out[104...]

```
tensor([[ -0.0000,       0.0000,       1.0000],
        [  0.0000,       1.0000,       0.0000],
        [ -1.0000,       0.0000,      -0.0000]])
```


In [105...]

```
meshes2 = pytorch3d.structures.Meshes(  
    verts=vertices @ relative_rotation,  
    faces=faces,  
    textures=textures,  
).to(device)  
  
# build a new mesh where each vertex p is mapped to p' = p R, with R = relative_rotation  
  
image = renderer(meshes2, cameras=cameras, lights=lights)  
plt.imshow(image[0].cpu().numpy())
```




Out[105...]

<matplotlib.image.AxesImage at 0x7f0fd5e870b0>




In [106...]


```
# Same output can be expected with the camera rotation  
  
cameras2 = pytorch3d.renderer.FoVPerspectiveCameras(  
    R=relative_rotation.unsqueeze(0),  
    T=[[0, 0, 3]],  
    device=device  
)  
  
image = renderer(meshes, cameras=cameras2, lights=lights)  
plt.imshow(image[0].cpu().numpy())
```



Out[106...]

<matplotlib.image.AxesImage at 0x7f0fd5ca7a40>





```
In [107]:  
    transform = cameras2.get_world_to_view_transform()  
    transform.get_matrix()
```


```
Out[107]: tensor([[[ -0.0000,  0.0000,  1.0000,  0.0000],  
                   [ 0.0000,  1.0000,  0.0000,  0.0000],  
                   [-1.0000,  0.0000, -0.0000,  0.0000],  
                   [ 0.0000,  0.0000,  3.0000,  1.0000]]],  
                  device='cuda:0')
```

```
In [108]:  
    plot_scene({  
        "Mesh1": {  
            "Mesh": meshes,  
            "Camera1": cameras,  
            "Camera2": cameras2,  
        },  
        "Mesh2": {  
            "Mesh": meshes2,  
            "Camera1": cameras,  
            "Camera2": cameras2,  
        },  
    }, ncols=2)
```

```
In [109]:  
    image = renderer(meshes2, cameras=cameras2, lights=lights)  
    plt.imshow(image[0].cpu().numpy())
```

```
Out[109]: <matplotlib.image.AxesImage at 0x7f0fd5527830>
```





In [110...]

```

num_views = 12
R, T = pytorch3d.renderer.look_at_view_transform(
    dist=3,
    elev=0,
    azim=np.linspace(-180, 180, num_views, endpoint=False),
)

# dist=3 places every camera on a circle of radius 3 around the origin
# elev=0 keeps them on the horizontal plane
# azim is an array of 12 angles from -180 to 180 degrees

print("R", R.shape)
print("T", T.shape)
many_cameras = pytorch3d.renderer.FoVPerspectiveCameras(
    R=R,
    T=T,
    device=device
)

# This builds a batch of 12 perspective cameras, one per view
images = renderer(meshes.extend(num_views), cameras=many_cameras, lights=li

```

R torch.Size([12, 3, 3])
T torch.Size([12, 3])


In [111...]

```

fig, axs = plt.subplots(3, 4)
axs = axs.flatten()
for i, image in enumerate(images):
    ax = axs[i]
    ax.imshow(image.cpu())
    ax.axis("off")

```





In [112...]

```
plot_scene({  
    "All Views": {  
        "Mesh": meshes,  
        "Cameras": many_cameras,  
    },  
})
```

Playing with Texture

In [113...]

```
texture_rgb = vertices.clone()  
texture_rgb = (texture_rgb - texture_rgb.min()) / (texture_rgb.max() - texture_rgb.min())  
texture_rgb /= texture_rgb.norm(dim=2, keepdim=True)  
textures_rainbow = pytorch3d.renderer.TexturesVertex(texture_rgb.to(device))  
meshes.textures = textures_rainbow
```





In [114...]

```
image = renderer(meshes, cameras=cameras, lights=lights)  
plt.imshow(image[0].cpu().numpy())
```

Out[114...]


```
<matplotlib.image.AxesImage at 0x7f0fd5afdeb0>
```





```
In [115...  
vertices, face_props, text_props = pytorch3d.io.load_obj("data/cow.obj")  
faces = face_props.verts_idx  
vertices = vertices.unsqueeze(0) # [1, V, 3]  
faces = faces.unsqueeze(0) # [1, F, 3]  
verts_uvs = text_props.verts_uvs  
faces_uvs = face_props.textures_idx  
  
texture_map = plt.imread("data/cow_texture.png")  
plt.imshow(texture_map)
```

Out[115... <matplotlib.image.AxesImage at 0x7f0fd41b1dc0>



```
In [116...  
textures_uv = pytorch3d.renderer.TexturesUV(  
    maps=torch.tensor([texture_map]),  
    faces_uvs=faces_uvs.unsqueeze(0),  
    verts_uvs=verts_uvs.unsqueeze(0),  
).to(device)
```


/tmp/ipython-input-2215026011_nv.2: 1: no pipenv

Creating a tensor from a list of numpy.ndarrays is extremely slow. Please consider converting the list to a single numpy.ndarray with `numpy.array()` before converting to a tensor. (Triggered internally at /pytorch/torch/csrc/utils/tensor_new.cpp:253.)

In [117...]

```
meshes.textures = textures_uv
image = renderer(meshes, cameras=cameras, lights=lights)
plt.imshow(image[0].cpu().numpy())
```

Out[117...]



Rendering Pointclouds

We will need:

- A point cloud
- A camera
- A point renderer

Setting up Point Cloud

In [118...]

```
coords = torch.randn(1000, 3)
rgb = torch.ones_like(coords) * torch.tensor([0.7, 0.7, 1])

print("Coords:", coords.shape)
print("RGB:", rgb.shape)
```

```
Coords: torch.Size([1000, 3])
RGB: torch.Size([1000, 3])
```

```
In [119...     pointcloud = pytorch3d.structures.Pointclouds(
                  points=coords.unsqueeze(0),
                  features=rgb.unsqueeze(0),
              ).to(device)
```


```
In [120...     raster_settings = pytorch3d.renderer.PointsRasterizationSettings(
                  image_size=512,
                  radius=0.01,
              )

              rasterizer = pytorch3d.renderer.PointsRasterizer(
                  cameras=cameras,
                  raster_settings=raster_settings,
              )

              renderer = pytorch3d.renderer.PointsRenderer(rasterizer=rasterizer,
                  compositor=pytorch3d.renderer.AlphaCompositor(background_color=(1, 1, 1),)
```

```
In [121...     image = renderer(pointcloud, cameras=cameras)
                  plt.imshow(image[0].cpu())
```

```
Out[121... <matplotlib.image.AxesImage at 0x7f0fc20153a0>
```



```
In [122...     plot_scene({
                  "Pointcloud": {
```

```

        "Pointcloud": pointcloud,
        "Camera": cameras,
    },
})

```

In [123...]

```

rgb = coords.clone()
rgb = (rgb - rgb.min()) / (rgb.max() - rgb.min())

pointcloud = pytorch3d.structures.Pointclouds(
    points=coords.unsqueeze(0),
    features=rgb.unsqueeze(0),
).to(device)

```

In [124...]


```

image = renderer(pointcloud, cameras=cameras)
plt.imshow(image[0].cpu())

```

Out[124...]

<matplotlib.image.AxesImage at 0x7f0fd5490f80>



Rendering Parametric Functions

Equation of a sphere with radius r with center (x_0, y_0, z_0) :

$$x = x_0 + r * \sin \theta \cos \phi$$

$$y = y_0 + r * \cos \theta$$

$$z = z_0 + r * \sin \theta \sin \phi$$

In [125...]

num_samples = 1000

```
num_samples = 1000
r = 1
x_0 = 0
y_0 = 0
z_0 = 0

phi = torch.linspace(0, np.pi, num_samples)
theta = torch.linspace(0, 2 * np.pi, num_samples)
Phi, Theta = torch.meshgrid(phi, theta, indexing="ij")

# Theta as polar and Phi as azimuth
print("Phi", Phi.shape)
print("Theta", Theta.shape)
```

```
Phi torch.Size([1000, 1000])
Theta torch.Size([1000, 1000])
```

```
In [126...]: x = x_0 + r * torch.sin(Theta) * torch.cos(Phi)
y = y_0 + r * torch.cos(Theta)
z = z_0 + r * torch.sin(Theta) * torch.sin(Phi)
```

```
In [127...]: points = torch.stack((x.flatten(), y.flatten(), z.flatten()), dim=1)
color = (points - points.min()) / (points.max() - points.min())

sphere_point_cloud = pytorch3d.structures.Pointclouds(
    points=[points], features=[color],
).to(device)
```

```
In [128...]: image = renderer(sphere_point_cloud, cameras=cameras)
plt.imshow(image[0].cpu())
```

```
Out[128...]: <matplotlib.image.AxesImage at 0x7f0fd63562d0>
```

